

Malloc

Fredrik Öberg
fobe@kth.se

Examiner: Prof. Johan Montelius

KTH Royal Institute of Technology, Sweden

January 7, 2021

Abstract

This paper examines how memory management works in a computer system and touches on topics like how allocation and deallocation of physical memory to running processes work and how those practices could be implemented. It also goes through what memory overhead is, and how it affects a computers memory. The paper shows that with som added complexity to a memory management algorithm a significant increase in memory utilage can be accomplished. This is shown through tests where both internal and external memory fragmentation becomes limited as well as a decrease in memory overhead after the implementation of the added complexity.

Contents

1 Introduction

This paper is written as a part of the examination of the course ID1206 - Operating Systems; itself part of the operations of KTH Royal Institute of Technology. The course provides knowledge of the principles of abstractions concerning computer hardware as well as virtualization of resources and timetabling of assignments and how those principles can be implemented; mainly as regards to program execution, memory management and persistent storage of data.

The purpose of this paper is to get a deeper understanding of how an operating system(OS) handles physical computer memory allocation and deallocation – a process by which computer programs and services are assigned physical memory. This is done by implementing different memory management algorithms and perform testing procedures on those algorithms to observe their behavior. The paper also touches on the topic of what can be done to increase both the efficiency in terms of memory usage as well as sheer performance.

2 Background

A process is an executing instance of a computer program. The process needs computer memory so it can perform its intended tasks and asks the computers OS for that amount of memory. This is accomplished by calling the OSs memory allocator through an application programming interface(API) the OS provides. The memory space is divided into blocks of different size and if the OS finds an available amount of memory satisfying the processes request, it marks that block of memory as taken and distributes it to the process. When the process is finished using its given memory block it needs to call the allocator once again through the OS API but this time telling it to free up the previously used memory block thereby making it available to other processes to use.

Each block of memory is assigned its specific header so the OS can keep track of the topology of the memory. The headers stores block information such as the size of the block, if it is free or not, how adjacent blocks look like etc. all depending on what type of memory manage algorithms are used by the memory allocator. One goal to ensure that the memory usage is efficient is to keep the size of the head as small as possible. This is especially important in computers with little memory since a large memory overhead used by the block heads then makes up a larger percentage of the complete usage of the memory i.e. less memory for the processes to do their work in. How a memory allocator manages the allocation and deallocation of memory also affects how the memory of a computer system behaves long term. An inefficient allocation algorithm gives larger blocks of memory to the processes than they need, thereby creating large internal fragmentation of the memory, since a large part of the allocated memory is not used. This situation can be somewhat managed by splitting up large blocks into smaller blocks with sizes more fit to the needs of the processes. An inefficient deallocation algorithm will create problems as well if it does not handle the freed-up memory in an efficient way. With an allocating

algorithm splitting up memory blocks into smaller sizes and a deallocator not merging adjacent free memory blocks a lot of external fragmentations will be created since the memory will be split up into smaller and smaller blocks of memory. This will create a lot of memory overhead since the memory used for block heads will increase which each new block created, and it will also create problems for future processes since there might not be blocks large enough for the OS to hand out which satisfies the process requests.

3 Set Up

To test out ways to solve the problems raised in the previous part, a basic allocating function was developed using the programming language C to check how different allocating algorithm works and what kind of behavior they bring to the memory allocating scheme. First was a memory block of 64 kilobytes(kB) allocated from the OS via the Unix system call `mmap` - a block from this point on called the arena - and by that simulating a system with that amount of memory available for the faux processes of the tests to use. A list of free blocks was created to keep track of how much memory there was available at a given time. The block headers used to store the block information was also created and these headers were in the initial allocating version 24 bytes large. They consisted of variables keeping track of the size and status of the current as well as the preceding block in the arena. Two variables were also implemented in the headers called `next` and `prev` for the free list to use. They enabled for a block who itself was on the list to keep track of what free block precedes as well as succeeds it on the free list. The boundaries of the previously created arena were set by two of these blocks so that the program could know in what frame the memory blocks could be located. The initial free head block thereby consisted initially of 65488 bytes, since 64kB minus 48 bytes for the boundary headers sums up to that amount.

The first allocation algorithm implemented had a way of splitting up free blocks into a size more fit to the current memory request in an attempt to limit internal fragmentation. It had on the other hand no way of merging free blocks which the next version of the allocation algorithm had. With some added complexity it checked if the blocks before and after it in the arena was also free and merged with them if that was the case creating a larger free block with only one header. The third and last version of the algorithm a smaller header than the initial implemented, consisting of only eight bytes worth of variables. The aim was to accomplish a smaller header overhead by adding some complexity to the algorithm without decreasing performance in an unreasonable way. This was enabled by removing the `next` and `prev` variables if a block was not in the free list.

The different memory management algorithms were tested by simulating several allocation and deallocation requests of variable sizes. The requests were between 8 and 1024 bytes in the large header version and 16 and 1024 bytes with the smaller one. The 8 bytes comes from that it is the smallest address

which can be used if the allocation is to be aligned with the memory of a 64-bit processor. The 16 bytes on the other hand comes from that the smaller header still needs to occupy a minimum of 24 bytes for the implemented free list to use. 16 bytes also aligns with the 64-bit address. Each request size was randomized with the bottom boundary was set at 8 and the top at 1024 bytes.

The testing involved using 1000 iterations of requests where the process started by allocating memory to 40 requests simulating a startup of a system. It thereafter had a 50 percent chance per iteration of making an allocating and/or deallocating request, simulating processes being finished as well as new being initiated.

Lastly a performance test was initiated where the performance of the small and the larger header versions was tested and compared to each other by allocating 1000 blocks of 16 bytes write to each of those blocks for 10000 iterations. This was done 100 times consecutively and the time it took for each one of those 100 write session were timed to see if the memory overhead of the different sized blocks had any effect on the computer's performance.

4 Results

The first graph (see Figure ??) is the result of the benchmark where no merging of contiguous free blocks were initiated. The graph shows a steady increase in the free list size though it seems beginning to pan out at the end. The graph also shows as a steady decrease in the average size of the free blocks available.




Figure 1: Free list and average free block size development without merging scheme

The second graph (see Figure ??) shows the result of the test where the merging algorithm had been implemented. It shows an increase in the free list size and decrease of the average block size until about 400 request iterations had been made. It then stays at a free list size of about 30 and an average free block size of slightly above 1000 bytes.

The third graph (see Figure ??) is the result of the test where the smaller block heads had been implemented along with the merging scheme. After about 300 request iterations the free list is quite stable at a size of slightly above 20 and the average block size is hovering around 2000 bytes.

The last graph (see Figure ??) shows the performance test difference between the different header type merging algorithms. It shows that the smaller head version stabilizes at about 6 milli seconds (mS) and the larger version at about 8 mS after an estimated 30 test executions.




Figure 2: Free list and average free block size development with merging scheme

5 Discussion

The result visualized in the first graphs clearly show that the size of the free list steadily increases as time passes. This is not surprising since for each request of a given size, either a block of exact that size is found in the free list, or a free block of a larger size found and split into two, creating a new block with its own header. Since the algorithm had no way of merging free blocks that means that the free list always increases in size for each deallocated block and as consequence results in free blocks with smaller and smaller sizes.




Figure 3: Free list and average free block size development with merging scheme and smaller block head

The second graph indicates that the merging algorithm has a free list size which is a sort of a probabilistic equilibrium. When the number of free blocks reaches above this equilibrium it is likely that they will get adjacent free blocks with which it can merge thereby decreasing the size free list. That the free list size and the average block size hovers around the values they do could be the result of the test implementation and those values could change with another test. However, it seems likely that those tests would show a similar behavior as this one, hovering around a steady value.

The result of the third test is very similar to the second but shows some increase in the memory usage. The result indicates that the free list size equilibrium in this test seem to be at around 20, or at least below what was seen in the test with the larger headers. The test also indicates that the average free block size is larger. These observations seem reasonable since a smaller memory overhead should result in larger block sizes overall as well as a smaller free list, at least when you test them the same way. That the difference is not larger when the head size of the smaller is a third of the size of the larger heads could be that the memory overhead contribution of the smaller heads becomes less if the average size of the requested block is large. In the tests the average request block should be at around 512 bytes with the previously mentioned request boundaries set in these tests, so the decrease in memory usage becomes a small fraction of that size. Even though the difference is relatively small it is still interesting to see that there at least is some difference between the tests with different header sizes.

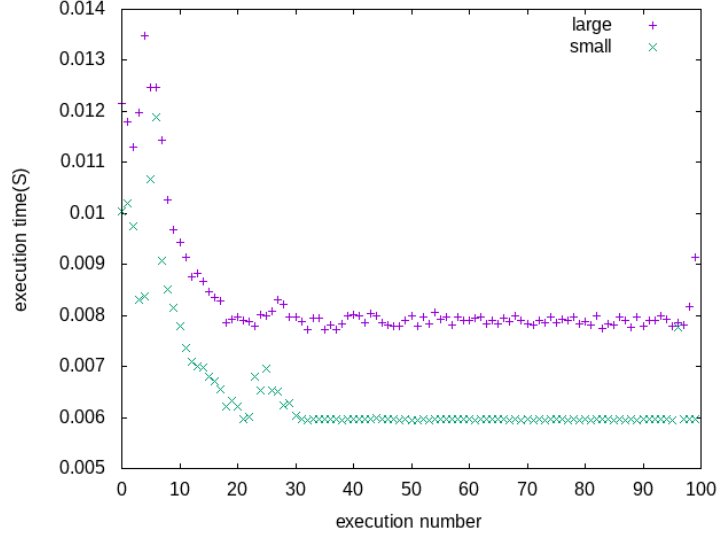


Figure 4: Execution time comparison between the small and large headers

The last graph (see Figure ??) clearly shows that we get a performance increase in having a smaller memory overhead, at least when the test is done where the average allocated block is as small as 16 bytes. At first the tests showed no major difference between using the different block heads, but with some optimization during code compilation the result is what can be seen in the graph. This could be that the difference is so insignificant (a time scale of mS) when the code is not optimized, and by removing some unnecessary parts of the code the difference shows more clearly.

The increase in performance probably comes from that the decrease in memory overhead from 24 to 8 per allocated block enables a total decrease in the allocated memory area from 40,000 to 24,000 bytes – a decrease in 40 percent. This makes it so that it is easier to store the blocks in the computers cache memory for faster access to the processor. The smaller block heads also result in a smaller number of bytes needed to be traversed which could contribute to the performance increase, but how much this contributes compared to the better use of cache memory these tests do not show.

6 Conclusions

It can be clearly stated that the memory allocating algorithm an operating system uses to allocate and deallocate the physical memory space affects the overall performance of a system. An effective algorithm for handling freed memory blocks by merging contiguous free blocks together decreases the memory overhead as well as memory fragmentation - both internal and external.

These improvements will of course come with the cost of more complex algorithms using up resources to accomplish this increase in memory usage. The process of finding the perfect balance in the trade-off between time versus space in computer resource management is something which has been an important part of computer resource management throughout the history of computers and will continue until the day of infinite computing resources - a day which most likely never will come.

The algorithms tested in this paper are similar to the ones used by modern OSs of today and the fact that modern computer work with memory sizes in the order of gigabytes instead of kilobytes used in this paper makes it unlikely that the performance increase from scaling down on header sizes will have a huge effect on modern computers overall performance - so those performance test results should not be overstated. But it should still be a goal to minimize memory overhead - just not at all cost. The process of merging contiguous free blocks however showed to have a big impact in these tests and should show similar behaviour in modern computers, so it is understandable that merging is a standardized practice. Finding even better solutions to the time-space trade off is a continuous process and it will be interesting to see what the future has to offer in the world of memory management.