

Green, green threads of home

Johan Montelius

HT2019

1 Introduction

This is an assignment where you will implement your own thread library. Instead of using the operating systems threads you will create your own scheduler and context handler. Before even starting to read this you should be up and running using regular threads, spin locks, conditional variables, monitors etc. You should also preferably have done a smaller exercise that shows you how we can work with *contexts*.

Note - the things you will do in this assignment are nothing that you would do in real life. Whenever you want to implement a multi-threaded program you would use the *pthread library*. We will however manage contexts explicitly to implement something that behaves similar to the pthread library. Why? - To better understand how threads work on the inside.

We will call our implementation *green* since they will be implemented in user space i.e. the operating system will not be involved in the scheduling of threads (it does perform the context switching for us so the threads are not all green).

2 Managing contexts

The library functions that we will use are: *getcontext()*, *makecontext()*, *setcontext()* and *swapcontext()*. Look up the man pages for these functions so that you get a basic understanding of what they do.

The functions that we want to implement are the following (arguments will be described later):

- *green_create()* : initializes a green thread
- *green_yield()* : suspends the current thread and selects a new thread for execution
- *green_join()* : the current thread is suspended waiting for a thread to terminate

Our handler needs to keep track of the currently *running* thread and a set of suspended threads *ready* for execution. It must of also keep track of threads that have terminated since the creator of the thread might want to call *green_join()* and should then be able to pick up the result.

We will try to mimic the *pthread* library so look up the definition of the equivalent functions: *pthread_create()*, *pthread_yield()* and *pthread_join()* (note - *green_join()* will take a pointer to a struct where *pthread_join()* takes a struct).

representing a thread

To follow the structure of the pthread library we will represent threads by a structure. This structure will hold all information that is needed for us to manage the threads; the threads manager will keep a very small state internally.

The structure should of course hold a pointer to the *context* of the thread. The context is used when we call *swapcontext()* or *setcontext()*. When we call the thread for the first time, we will call a *function* given an *argument*. When a thread is suspended we will add it to a linked list so we include a *next pointer* in the structure itself.

We also need to keep track of a thread that is waiting to *join* the thread to terminate so we also keep a pointer for this and a pointer to the final result- The *zombie* status field will indicate if the thread has terminated or not.

In a file *green.h* include the following:

```
#include <ucontext.h>

typedef struct green_t {
    ucontext_t *context;
    void *(*fun)(void*);
    void *arg;
    struct green_t *next;
    struct green_t *join;
    void *retval;
    int zombie;
} green_t;

int green_create(green_t *thread, void *(*fun)(void*), void *arg);
int green_yield();
int green_join(green_t *thread, void** val);
```

internal state

The internal state of the scheduler is very small. We will only keep track of two things: the *running* thread and the *ready queue*. We will also need one global context that will be used to store the main context i.e. the context of the initial running process.

In a file *green.c* include the following:

```

#include <stdlib.h>
#include <ucontext.h>
#include <assert.h>
#include "green.h"

#define FALSE 0
#define TRUE 1

#define STACK_SIZE 4096

static ucontext_t main_ctx = {0};
static green_t main_green = {&main_ctx, NULL, ... FALSE};

static green_t *running = &main_green;

```

We allocate a global green thread *main_green* that holds a pointer to the main context and otherwise initialized with null pointers and a false zombie flag.

If you read the man pages for *makecontext()* and *swapcontext()* you realize that we will have to initiate a contexts before we can use it and this can not be done at compile time. We could provide a function that the user needs to call to initialize the manager, or a in the first call to *green_create()* detect that the main context needs to be initialized; we will however use another trick - provide a function that is called when the program is loaded.

In the file *green.c* add the following definition.

```

static void init() __attribute__((constructor));

void init() {
    getcontext(&main_ctx);
}

```

The *init()* function will initialize the *main_ctx* so when we call the scheduling function for the first time the *running* thread will be properly initialized.

2.1 create a green thread

A new green thread is created in a two stage process. First the user will call *green_create()* and provide: an uninitialized *green_t* structure, the *function* the thread should execute and pointer to its *arguments*. We will create new context, attach it to the thread structure and add this thread to the *ready queue*.

When this thread is scheduled it should call the function but in order to do this we set it to call the function *green_thread()*. This function is responsible for calling the function provided by the user.

```

int green_create(green_t *new, void *(*fun)(void*), void *arg) {

    ucontext_t *cntx = (ucontext_t *)malloc(sizeof(ucontext_t));
    getcontext(cntx);

    void *stack = malloc(STACK_SIZE);

    cntx->uc_stack.ss_sp = stack;
    cntx->uc_stack.ss_size = STACK_SIZE;
    makecontext(cntx, green_thread, 0);

    new->context = cntx;
    new->fun = fun;
    new->arg = arg;
    new->next = NULL;
    new->join = NULL;
    new->retval = NULL;
    new->zombie = FALSE;

    // add new to the ready queue
    :

    return 0;
}

```

It is up to you to implement how the ready queue is managed.

Now let's take a look at the *green_thread()* function. This function will do two things: start the execution of the real function and, when after returning from the call, terminate the thread.

```

void green_thread() {
    green_t *this = running;

    void *result = (*this->fun)(this->arg);

    // place waiting (joining) thread in ready queue
    :
    // save result of execution
    :
    // we're a zombie
    :
    // find the next thread to run
    :
    running = next;
    setcontext(next->context);
}

```

```
| }
```

The tricky part is what to do when the called function returns. We (that is you) should check if there is a thread waiting for its termination, and if so place it in the ready queue; the thread is now a zombie process. The result of the execution should be saved to allow the waiting thread to collect the result.

There should be a thread in the ready queue so we select the first and schedule it for execution.

2.2 yield the execution

In the initial implementation, scheduling is only done when a thread voluntarily call the *green_yield()* function. This function will simply put the running thread last in the ready queue and then select the first thread from the queue as the next thread to run.

```
| int green_yield() {  
|     green_t * susp = running;  
|     // add susp to ready queue  
|     :  
|     // select the next thread for execution  
|     :  
|     running = next;  
|     swapcontext(susp->context, next->context);  
|     return 0;  
| }
```

The call to *swapcontext()* will do the context switch for us. It will save the current state in *susp->context* and continue execution from *next->context*. Note that when the suspended thread is scheduled, it will continue the execution from exactly this point (read that sentence again, it will be important).

2.3 the join operation

The join operation will wait for a thread to terminate. We therefore add the thread to the *join* field and select another thread for execution. If the thread has already terminated we can of course continue as if nothing happened. In either case we will of course pick up the returned value from the terminating thread. We will also free the memory allocated by the zombie thread, it is now dead.

```
| int green_join(green_t *thread, void **res) {  
|  
|     if(!thread->zombie) {  
|         green_t *susp = running;  
|         // add as joining thread
```

```

        :
        //select the next thread for execution
        :
        running = next;
        swapcontext(susp->context, next->context);
    }
    // collect result
    :
    // free context
    :
    return 0;
}

```

What will happen if several threads wait for the same thread? If you read the man pages for *pthread_join()* you will see that they say that the behavior is undefined. This is a reasonable decision and makes life easier for us. Let's adopt the strategy and only allow one waiting thread.

2.4 a small test

If you have completed the code above and implemented a ready queue, you should be able to run a small test. Separate the test program from your implementation and include the header file *green.h*.

```

#include <stdio.h>
#include "green.h"

void *test(void *arg) {
    int i = *(int*)arg;
    int loop = 4;
    while(loop > 0) {
        printf("thread %d: %d\n", i, loop);
        loop--;
        green_yield();
    }
}

int main() {
    green_t g0, g1;
    int a0 = 0;
    int a1 = 1;
    green_create(&g0, test, &a0);
    green_create(&g1, test, &a1);

    green_join(&g0, NULL);
}

```

```

    green_join(&g1, NULL);
    printf("done\n");
    return 0;
}

```

3 Suspending on a condition

Now for the next task: the implementation of conditional variables. These should work as the conditional variables in the pthread library. However, we do not have any mutex structures that can be locked so our implementation is simpler.

You should stop and wonder why we have not implemented any locking functionality. Is it not very dangerous to run multi-threaded programs without locks?

You should implement the following functionality:

- *void green_cond_init(green_cond_t*)*: initialize a green condition variable
- *void green_cond_wait(green_cond_t*)*: suspend the current thread on the condition
- *void green_cond_signal(green_cond_t*)*: move the first suspended thread to the ready queue

You need to define a data structure *green_cond_t* that can hold a number of suspended threads. The implementation of the functions should then be quite simple. Draw some pictures that describes what the operations should do before you start implementing.

When you think you have it you could try something like this (don't forget to initialize the conditional variable):

```

int flag = 0;
green_cond_t cond;

void *test(void *arg) {
    int id = *(int*)arg;
    int loop = 4;
    while(loop > 0) {
        if(flag == id) {
            printf("thread %d: %d\n", id, loop);
            loop--;
            flag = (id + 1) % 2;
            green_cond_signal(&cond);
        } else {

```

```

    green_cond_wait(&cond);
}
}
}

```

4 Adding a timer interrupt

So far we have relied on the threads themselves to either yield the execution or suspend on a conditional variable, before we schedule a new thread for execution. This is fine, and for sure makes things easier, but we might want to allow several threads to execute concurrently. We therefore introduce a timer driven scheduling event.

Note - `printf()` is not `asynch_signal_safe` which means that havoc will follow if we have a timer interrupt while printing. You need to use the system call `write()` if you want to log what is happening.

A timer will be set to send the process a signal with regular intervals. When we receive a signal we will suspend the currently running thread and schedule the next one in the run queue. This is exactly what the `green_yield()` function does.

In the beginning of *green.c*:

```

#include <signal.h>
#include <sys/time.h>

#define PERIOD 100

static sigset_t block;

void timer_handler(int);

```

Now in the `init()` function we initialize the timer. We initialize the *block* to hold the mask of the *SIGVTALRM*, set the handler to our `timer_handler()` function and associate the signal to the action handler. We then set the timer interval and delay (*value*) to our *PERIOD* and start the timer.

```

sigemptyset(&block);
sigaddset(&block, SIGVTALRM);

struct sigaction act = {0};
struct timeval interval;
struct itimerval period;

act.sa_handler = timer_handler;
assert(sigaction(SIGVTALRM, &act, NULL) == 0);

```



```

interval.tv_sec = 0;
interval.tv_usec = PERIOD;
period.it_interval = interval;
period.it_value = interval;
setitimer(ITIMER_VIRTUAL, &period, NULL);

```

When the timer expires the handler will be called and its time to schedule the next thread.

```

void timer_handler(int sig) {
    green_t * susp = running;

    // add the running to the ready queue

    // find the next thread for execution
    running = next;
    swapcontext(susp->context, next->context);
}

```

If you complete the code above it will actually work ... almost. You could test it for a while before you run into a strange segmentation fault and when you do, you will have a very hard time finding the bug.

The thing that will eventually happen is that we will have a timer interrupt when we're in one of the functions that manipulate the state of the green threads. Imaging what could happen if we are in the middle of a yield operation and change the run queue. We need to prevent these interrupts when we change the state.

Fortunately for us, we are not the only ones with this problem so we have a simple way to block and unblock these interrupts:

```

sigprocmask(SIG_BLOCK, &block, NULL);
:
:
:
sigprocmask(SIG_UNBLOCK, &block, NULL);

```

5 A mutex lock

When you have your timer interrupts working, write a small test program that shows that it works. Also write a program that shows a situation where our threads library falls short.

Since a thread now can be interrupted at any point in the execution we will have a problem when we update shared data structures. A simple example where two threads read and increment a shared counter will lead to very unpredictable behavior. We need a way to synchronize our threads and a mutex construct would do the trick.

We will need a structure to represent a mutex and since we will have threads suspended on the lock we let it hold a list of suspended threads. In the *green.h* file we add the following:

```
typedef struct green_mutex_t {
    volatile int taken;
    // handle the list
} green_mutex_t;

int green_mutex_init(green_mutex_t *mutex);
int green_mutex_lock(green_mutex_t *mutex);
int green_mutex_unlock(green_mutex_t *mutex);
```

The function *green_mutex_init()* is trivial since all we have to do is initialize the fields:

```
int green_mutex_init(green_mutex_t *mutex) {
    mutex->taken = FALSE;
    // initialize fields
}
```

The function that tries to take the lock will, if the lock is taken, look very similar to the yield procedure. We will use a method called “pass the baton” - a thread that holds the lock will give the lock to the next thread in line. If we wake up after having suspended on the mutex we know that the lock is ours.

```
int green_mutex_lock(green_mutex_t *mutex) {
    // block timer interrupt

    green_t *susp = running;
    if(mutex->taken) {
        // suspend the running thread
        :
        // find the next thread
        running = next;
        swapcontext(susp->context, next->context);
    } else {
        // take the lock
        :
    }
    // unblock
    return 0;
}
```

The unlock function is very similar to the signal operation. If there is a thread waiting on the lock we do not release the lock but pass it over to the suspended thread.

```

int green_mutex_unlock(green_mutex_t *mutex) {
    // block timer interrupt
    :
    if(mutex->susp != NULL) {
        // move suspended thread to ready queue
    } else {
        // release lock
    }
    // unblock
    return 0;
}

```

Complete the code and write a small program that shows that it works.

6 The final touch

Take a look the procedure below, where two threads take turn changing a flag. The flag is protected by a mutex but we use a conditional variable to signal that the flag has changed. Will this work? What will happen if we release the lock and have a timer interrupt before calling `green_cond_wait()`?

```

while(loop > 0) {
    green_mutex_lock(&mutex);
    while(flag != id) {
        green_mutex_unlock(&mutex);
        green_cond_wait(&cond);
        green_mutex_lock(&mutex);
    }
    flag = (id + 1) % 2;
    green_cond_signal(&cond);
    green_mutex_unlock(&mutex);
    loop--;
}

```

We need, as in the pthread library, a function that suspends on a conditional variable and releases a lock in one atomic operation. When the function returns, the lock should be held. We therefore change the function `green_cond_wait()` to also take a mutex as an argument.

```

int green_cond_wait(green_cond_t *cond, green_mutex_t *mutex) {
    // block timer interrupt
    :
    // suspend the running thread on condition
    :
    :

```

```

    if(mutex != NULL) {
        // release the lock if we have a mutex
        :
        // move suspended thread to ready queue
        :
    }
    // schedule the next thread
    :
    running = next;
    swapcontext(susp->context, next->context);

    if(mutex != NULL) {
        // try to take the lock
        if(mutex->taken) {
            // bad luck, suspend
            :
        } else {
            // take the lock
            mutex->taken = TRUE;
        }
    }
    // unblock
    :
    return 0;
}

```

Rewrite your test program and use the atomic conditional wait function. You should now be able to have a producer and consumer synchronize their actions using conditional variables and mutex locks.

7 Summary

A threads library is not rocket science but it takes some time before you get it right. The problem is that when it fails it is very hard to figure out what went wrong. Does your implementation work correctly? A single threaded program is easier to debug since it follows the same execution path every time you run it. In our implementation things were predictable up until we introduced the timer interrupts. How do we know that we have covered all corner situations? Is extensive testing the only tool we have? If your threads library was controlling the threads of a web browser it might not matter very much, what if it was used in an airplane control system - would you sleep well tonight?